

SkyBridge Compass: A Crypto-Agile P2P Handshake Framework for Remote Desktop and File Transfer with Auditable Downgrade Policy and Post-Quantum Readiness

Zi'ang Li, Peng Liu

Abstract—Peer-to-peer remote desktop and file transfer systems increasingly operate across heterogeneous devices, untrusted networks, and long lifecycles in which cryptographic assumptions may expire faster than products. Yet “post-quantum readiness” is often treated as a one-off algorithm swap rather than a migration problem with downgrade resistance, auditability, and reproducibility requirements. This paper presents an auditable cryptographic migration contract, instantiated in SkyBridge Compass, that explicitly separates (i) negotiated cryptographic suites, (ii) provider backends, and (iii) a policy layer with explicit fallback semantics and a minimal evidence set for replaying downgrade decisions. Under Default Mode, classic establishment is never triggered by timeout or network-derived failures; any allowed classic fallback is local-only (pre-send), rate-limited, and emits replayable downgrade evidence. We use an actor-isolated handshake state machine with deterministic outcomes and structured security-event emission, providing protocol-level audit signals that make downgrade decisions and failure semantics explicit and replayable. Signals are linked to transcript-derived identifiers for in-session tamper-evident ordering under a non-compromised runtime assumption; long-term retention and remote attestation of logs are out of scope. We evaluate a reproducible *core gate* on macOS 26.x with three configurations: classic (X25519 + Ed25519), PQC via liboqs (ML-KEM-768 + ML-DSA-65), and liboqs PQC v2 FS (ML-KEM-768-FS + ML-DSA-65). Apple CryptoKit PQC/X-Wing are reported as a separate *contrast* dataset and do not participate in hard pass/fail gates. To avoid cross-domain confounding, we evaluate source-quality, paper-integrity, and release-readiness gates independently with traceable pass/fail evidence in each domain. To prevent upgrade-time moving-goalpost effects, this track freezes wire formats, suite/wire IDs, canonical vectors, the artifact date (2026-01-23), threat model, and strict/prefer/classic policy semantics (including fallback-deny behavior). Cross-platform uplift is validated with a decision-parity matrix that fixes policy mode, peer visibility, local provider availability, compile-time flags, trust-store state, and bootstrap/rekey state, then requires equivalent suite choice, fallback allow/deny, event semantics, and user-facing error categories across macOS and iOS. We do not claim perfect forward secrecy for the KEM-based PQC suites in v1; post-quantum confidentiality holds only while long-term KEM private keys remain uncompromised. End-to-end handshake completion latency (including event emission overhead) is 2.232 ms (p95 2.417 ms) for classic, 2.743 ms (p95 3.404 ms) for liboqs PQC, and 3.483 ms (p95 4.184 ms) for liboqs PQC v2 FS under the core benchmark controls (N=1000 per configuration). Apple contrast runs are collected separately (N=5000) and reported in the supplementary Apple/X-Wing contrast report. The handshake payload size (MessageA + MessageB + 2×Finished) grows from 687 B (classic) to 12002 B (liboqs PQC) and 12077 B (liboqs PQC v2 FS) in the core gate. Apple PQC/X-Wing payloads are reported as contrast-only data in the supplementary Apple/X-Wing contrast report. Beyond performance, we validate downgrade policies, failure semantics, and audit-signal fidelity under fault injection, and we release scripts and artifacts that reproduce all primary figures and tables. Our portability claim is mechanism-level portability of a provider/suite policy contract (not cross-platform performance equivalence); reported latency/overhead values are deployment-instance measurements.

Index Terms—post-quantum cryptography, crypto agility, downgrade resistance, secure handshake, peer-to-peer systems, remote desktop, auditability, reproducibility, platform deployment.

1 INTRODUCTION

Peer-to-peer (P2P) remote desktop and file transfer systems must establish secure sessions under adversarial network conditions, device churn, and multi-year lifecycles. These systems face a practical tension: users expect seamless pairing across heterogeneous stacks, while defenders require that cryptographic negotiation remains robust against downgrade and implementation-level failure modes. Post-quantum cryptography (PQC) increases this tension. Larger messages, higher computational cost, and uneven platform support make fallback tempting—yet fallback is exactly where

silent downgrade and unverifiable behavior tend to hide.

Existing secure-channel designs and frameworks provide mature cryptographic building blocks, but they do not, by default, solve the migration problem an implementer actually faces: (i) selecting among classic/PQC/hybrid suites across provider/suite strata in heterogeneous runtimes, (ii) ensuring that policy (e.g., “require PQC”) is enforced at the entry points where callers cannot accidentally violate it, (iii) targeting deterministic failure semantics under timeouts, reorder/duplication, and malformed inputs, and (iv) emitting security-relevant telemetry that can be audited without relying on ad hoc or excessive logging. In practice, many systems either hard-code a single suite, or allow “compatibility” to creep into negotiation in ways that

- Zi'ang Li is with Independent Researcher, Tianjin, China. E-mail: 2403871950@qq.com.
- Peng Liu is with Independent Researcher, Qiqihar, China.

TABLE 1
Migration requirements and default baseline coverage in Introduction scope.

Req.	Widely deployed defaults	SkyBridge contract in this paper (TLS/QUIC/DTLS/Noise)
(i) Suite/provider migration	Protocol negotiation primitives exist (Noise patterns may be fixed), but runtime provider migration policy is mostly external to the protocol	Explicit separation of suite stratum with capability-ordered selection
(ii) Policy non-bypass at entry points	Typically enforced in caller glue; policy/suite mismatch handling is implementation-specific	Policy-bound handshake sessions plus compile-time precondition checks (Supplementary, Regression Testing)
(iii) Deterministic fault semantics	Alert/pattern outcomes exist, but fallback semantics are often application-defined	Actor-isolated one-shot state machine with timeout-blocked fallback and deterministic terminal templates
(iv) Auditable security telemetry	Downgrade prevention exists, but event schema and evidence sufficiency are externalized	D3 minimal evidence set + typed events + transcript-anchored in-session hash chain

are hard to reason about and harder to verify. For this paper, “locally generated” fallback causes are narrowly scoped: pre-send capability/provider instantiation failures for already paired peers (pinned trust record); timeout and unauthenticated network-derived failures are never fallback-eligible.

To make the migration gap explicit in reviewer-visible form, Table 1 maps the four requirements above to what widely deployed handshakes provide by default versus what this work adds as a system-level contract.

This paper addresses these gaps with an auditable migration contract, instantiated in SkyBridge Compass, for staged PQC migration and downgrade-auditable negotiation, with Apple platforms as a deployment instance.

Distinction from protocol-plus-policy wrappers: In widely deployed stacks, the causes and semantics of retry/fallback are typically pushed into application glue, making downgrade decisions hard to specify, prove, or audit post hoc. Our contract makes migration semantics part of the security surface: (D2) timeout and other network-derived failures are never fallback-eligible, so an active adversary, under this default-policy gate and threat model, cannot create a classic downgrade edge by suppression; (D3) any policy-allowed downgrade action emits a minimal replayable evidence tuple sufficient to re-derive “what happened and why” under the pinned trust record and policy mode. The policy-bound, deterministic state machine ensures evidence corresponds to the actual decision path and is validated under fault injection and timeout suppression (Table 5).

The core idea is to treat crypto agility as a layered contract: negotiated suites define the protocol surface; provider backends implement suites using platform-native or portable cryptography; and a policy layer governs fallback and produces downgrade-specific structured evidence only when fallback action is taken (other exceptional states emit typed events but are not labeled as downgrade). Concretely, policy consumes peer-auth result, local capability/provider errors, offered suites, and timer/attempt state, and outputs allowed suites, retry plan, and mandatory event set. We complement this design with an actor-isolated handshake driver that enforces one-shot completion, timeouts, and sensitive-material zeroization, yielding deterministic outcomes under the modeled transport faults. The portability claim here is mechanism-level portability of the migration contract (provider/suite separation, policy semantics, and verification workflow) across hetero-

geneous stacks, not cross-platform latency parity. Operationally, the strict policy now has two explicit realizations: *Compliance Mode* (no classic path at all) and *Bootstrap-Assisted Mode* (one-time classic control channel only to refresh paired KEM identity keys, followed by mandatory PQC rekey before business traffic is released; control messages are type-restricted to KEM-refresh, and business message types are rejected before key release). We separate mechanism-level claims (policy-gated downgrade behavior, deterministic failure semantics, and auditability) from deployment-level claims (absolute latency/throughput), and evaluate each layer with corresponding evidence.

Verifiable Security Guarantees. To reduce ambiguity in reviewer-facing claims, we state the v1 guarantees as a small set of *verifiable* properties under the threat model in Section 5; the Supplementary Materials provide additional guarantees and artifact-level evidence pointers. Here “deterministic outcomes” means that for a fixed input transcript, policy mode, and fault class, the terminal outcome and required event template are unique.

- **Negotiation integrity:** the Responder selects only a suite offered by the Initiator and with a matching key share; unknown suite IDs in `keyShares[]` are rejected before signature verification (parse-stage), and unknown `selectedSuite` values are rejected during verification, so suite-ID handling does not become an oracle surface. *Evidence:* Fig. 1; Table 3.
- **Transcript binding:** `sigB` commits to `MessageA` and the chosen suite (and policy), so any tampering of `supportedSuites[]`, `keyShares[]`, or `policy` fails verification. Policy and suite lists are canonically encoded (v1 deterministic; v2 TLV canonical) to avoid equivocation. *Evidence:* Table 3.
- **Downgrade resistance and audibility (policy-defined):** strict policy is explicit and auditable in two forms: *Compliance Mode* forbids classic establishment entirely; *Bootstrap-Assisted Mode* permits a one-time classic control channel only for paired KEM-key recovery and requires PQC rekey before business traffic; this control channel is type-restricted to KEM-refresh messages. The default policy permits classic fallback only for a local-only whitelist of capability/provider failures (timeouts and unauthenticated network-derived failures cannot trigger default-mode fallback). Representative whitelist classes include local provider/API unavailability and local pre-send suite-selection failures (e.g., missing paired KEM keyshare). Every accepted fallback emits a structured `cryptoDowngrade` event. Non-fallback failures emit typed events (e.g., `handshakeFailed`, `identityMismatch`, `signatureVerificationFailed`) and are not labeled downgrade. *Evidence:* Fig. 4; Fig. 5; Table 7; Fig. 7.

Representative minimal event typings are summarized in Supplementary Table S37 (S37).

We additionally provide explicit key confirmation via two Finished frames (Fig. 1) and document v1 non-claims and limitations in Section 8. This aligns with standard AKE key-confirmation notions while matching our `MessageA/MessageB/two-Finished` structure (both peers confirm transcript-derived keys before business data release).

1.1 Contributions

Contributions:

- 1) **C1: Formalized auditable migration contract.** We formalize an auditable migration contract spanning four mechanism-level requirements (suite/provider migration, policy non-bypass, deterministic fault semantics, auditable evidence) and define four experiment families (authentication, negotiation, downgrade,

hybrid forward secrecy). We state theorem/proposition guarantees under standard assumptions (KEM IND-CCA, SIG EUF-CMA, HKDF PRF) and provide a machine-checkable Tamarin artifact aligned with implementation-level claims. *Evidence:* Section 5; Table 3; `formal/skybridge_v2.spthy`.

- 2) **C2: Protocol v2 upgrade with backward-compatible bridge.** We introduce a forward-secrecy-oriented v2 suite that composes a static KEM secret with an additional ephemeral contribution while retaining v1 interop through policy-controlled bridge negotiation. v1 remains default for bridge compatibility with provisioned KEM identity keys; v2 is an opt-in upgrade (not full PFS). *Evidence:* Fig. 1; Supplementary Table S22; `HandshakeV2PFSTests`; `ProtocolSignatureRegressionTests`.
- 3) **C3: Reproducible implementation and verification workflow.** We ship an actor-isolated implementation that emits protocol-level audit signals and maintains transcript-anchored, in-session tamper-evident ordering under a non-compromised runtime assumption; long-term retention and remote attestation depend on the operating environment and are out of scope. We also release artifact pipelines reporting distributional metrics (p50/p95/p99, CI/effect-size) rather than point estimates only. *Evidence:* Fig. 4; Table 9; `AuditTrailTests`; `HandshakeBenchmarkTests`.
- 4) **C4: Policy semantics for strict/default downgrade control.** We expose strict policy in two explicit modes (Compliance and Bootstrap-Assisted), define a local-only whitelist for default fallback, and require downgrade-specific evidence only when fallback action is taken; non-fallback failures are typed but not labeled downgrade. *Evidence:* Fig. 4; Fig. 5; Table 7; Section 6.

1.2 Paper Organization

Section 2 reviews related work on P2P pairing, negotiation, and PQC migration. Section 3 details system architecture and protocol design. Section 4 presents the handshake state machine. Section 5 presents the security model and guarantees. Section 6 describes implementation highlights and platform instantiation. Section 7 outlines evaluation methodology and metrics. Section 8 discusses limitations and future work, and Section 9 concludes. Supplementary Materials provide full wire-format definitions, implementation details, and extended tables.

2 RELATED WORK

We use the following terms throughout: *crypto agility* (ability to migrate primitives without redesign), *downgrade resistance* (prevent or detect forced fallback), *suite negotiation* (explicit, authenticated selection of handshake KEM/SIG/AEAD/KDF), *PQC readiness* (ability to deploy PQC alongside classic suites), and *auditability* (deterministic, minimal event traces for security decisions). Our threat model assumes an active network attacker attempting to induce fallback or obscure failures; therefore we focus on mechanisms that make negotiation and failures explicit and auditable.

2.1 Peer-to-Peer Device Pairing

Pairing protocols such as Bluetooth SSP [1], Continuity [2], and PAKEs like SPAKE2+ [3] establish initial trust but stop short of post-pairing migration, downgrade policy, or auditability. SkyBridge Compass begins *after* pairing and treats pinned identity keys as inputs rather than outcomes.

2.2 Cryptographic Agility and Negotiation

TLS 1.3 and QUIC bind negotiated parameters into transcripts and transport parameters to prevent silent downgrade [4], [5], but they do not define an auditable, system-level fallback policy across heterogeneous providers. We operationalize migration guidance [6], [7] with explicit policy gates and event emission that make downgrade decisions observable and testable.

2.3 Handshake Frameworks and Noise-Style Patterns

Noise provides a compact framework for describing authenticated key exchange patterns with explicit transcript binding and identity key usage [8]. While SkyBridge Compass does not implement Noise directly, our MessageA/MessageB/Finished exchange targets similar transcript-binding goals while adding explicit suite negotiation and auditable policy outcomes.

2.4 Post-Quantum Cryptography Deployment

PQC standardization and early deployments highlight size and performance costs, along with pressure for hybrid deployments [9], [10], [11], [28], [29], [12], [30], [31], [33], [13], [14], [15]. We focus on making downgrade policy and auditability explicit across providers, treating Apple’s CryptoKit PQC [16], [23] as a deployment instance rather than a protocol dependency.

2.5 Secure State Machine Design

State machine vulnerabilities have been a persistent source of security bugs in protocol implementations [17], [18]. Actor-based concurrency models, as implemented in Swift’s actor system [19], provide compile-time guarantees against data races. We extend this model with deterministic completion, timeout handling, and audit events for cryptographic handshakes.

2.6 Baseline Comparison with Widely Deployed Handshakes

Table 2 contrasts SkyBridge Compass with widely deployed handshake frameworks on explicit negotiation, downgrade auditability, and failure semantics, and reports message count, wire size, and latency where available.

Baseline protocols expose transcript integrity but do not specify system-level retry causes or replayable downgrade evidence. Detailed baseline methodology (time-to-ready endpoint, warmup policy, ALPN setup, kickoff payload settings, and loopback pcap accounting) and full baseline tables are provided in the Supplementary Materials (notably Table S24).

Importantly, baseline protocols’ downgrade protections target negotiated parameter integrity within a single handshake. They do not specify when an implementation should *retry* with weaker cryptography after local failures, nor do they require a minimal, replayable evidence set for such a decision; in practice, these behaviors are pushed into application glue and operational logging. Our migration contract makes the retry/downgrade boundary explicit (Definitions D2–D3) and audit-visible, and we validate it under active suppression (e.g., Fig. 5).

TABLE 2

Baseline comparison with widely deployed handshakes. Message counts and negotiation semantics are summarized from the TLS/QUIC/DTLS RFCs and Noise specification [4], [5], [8], [20]. Loopback latencies and wire sizes are measured with BaselineBenchRunner (release), N=1000; wire sizes report loopback p50/p95 (KiB) and are cert-dependent for TLS/QUIC/DTLS. SkyBridge loopback wire size includes framing/transport overhead (and may include SBP1 handshake padding); payload-only handshake bytes are reported in Supplementary Table S14. Noise wire size is pattern-dependent and reflects the XX pattern used in this baseline.

Protocol	Suite nego. explicit	Auditable downgrade policy	Failure semantics	Msgs	Wire size (p50/p95, KiB)	p50/p95 latency
TLS 1.3 [4]	Yes (cipher suites, groups)	Downgrade prevention only; audit external	Alerts defined; audit external	4 (1 RTT)	8.6/8.7	8.53/14.01 ms
QUIC [5]	Yes (TLS 1.3)	Same as TLS 1.3	Same as TLS 1.3	4 (1 RTT)	7.6/14.1	0.34/11.37 ms
WebRTC DTLS [20]	Yes (DTLS 1.2 ciphers)	Downgrade prevention only; audit external	Alerts defined; audit external	4 (1 RTT)	3.0/3.1	8.18/11.79 ms
Noise XX [8]	No (fixed pattern)	No policy layer	Pattern-defined; audit external	3	2.5/2.6	0.44/0.58 ms
					4.8/4.9 (classic)	1.44/1.57 ms (classic)
					36.9/43.6 (liboqs)	1.95/2.76 ms (liboqs)
SkyBridge Compass	Yes (suite + policy bound)	Yes (explicit policy + events)	Deterministic + evented	4	18.6/18.8 (CryptoKit)	5.91/7.53 ms (CryptoKit)

2.7 Novelty Boundary and Incremental Delta

To reduce novelty ambiguity, we separate reused protocol primitives from the system-level delta claimed in this paper. We do not claim novelty in authenticated transcript binding, suite negotiation primitives, or baseline actor isolation by itself. The claimed novelty is the *auditable migration contract*:

- **Explicit suite/provider separation** (the contract binds negotiation to a declared backend, not an implicit platform API default).
- **Policy non-bypass at entry points** (policy-bound session API + compile/runtime guards block ad hoc fallback in caller glue).
- **Narrow local-only fallback eligibility** (paired trust record, pre-send; timeout/network-derived causes are fallback-ineligible by definition) (D2).
- **Downgrade-specific audit semantics** (replayable evidence fields + deterministic event typing distinguish downgrade vs failure) (D3).
- **Explicit strictness realizations** (Compliance vs Bootstrap-Assisted strict recovery with mandatory PQC rekey gate before business traffic).

Supplementary Table S1 provides the full reviewer-facing prior-art vs. delta matrix and evidence anchors. *Ablation summary*: Supplementary Table S9 provides an artifact-backed ablation-by-property map for the contract mechanisms above.

2.8 PQC Migration Design Space

To make migration-risk boundaries explicit, we classify designs on two axes: (1) fallback trigger origin (*network-derived* vs *local-only* vs *never fallback*), and (2) evidence strength for operator-side verification (*none/external log* vs *structured replayable evidence*).

Representative points in this design space:

- **TLS/QUIC deployment defaults** [4], [5]: strong handshake transcript integrity, but system-level retry and fallback causes and evidence are implementation-defined outside the protocol.
- **Signal PQXDH** [10]: hybrid key agreement at the application layer, but system-level fallback eligibility and replayable downgrade evidence are outside the protocol.

- **SkyBridge migration contract (this work)**: local-only fallback whitelist with timeout/network-derived blacklist, plus structured replayable downgrade evidence and deterministic event typing.

Supplementary Table S2 provides the full matrix.

3 SYSTEM ARCHITECTURE

3.1 Overview

SkyBridge Compass implements a layered architecture separating concerns across four primary domains: discovery, cryptographic operations, handshake management, and session transport. Fig. 2 illustrates the high-level component relationships and trust boundaries.

3.2 Threat Model

We model an active network attacker in the Dolev-Yao style with the following **capability set**:

Network Capabilities (Attacker-Controlled):

- **Drop / delay / reorder / duplicate** packets on the discovery channel
- **Modify/inject** arbitrary bytes in MessageA, MessageB, or Finished frames
- **Replay** prior handshake messages across sessions
- **Spoof capability or policy claims** to induce weaker negotiation
- **Force negotiation failure** by corrupting suites, key shares, or signatures

Defense Mapping (Protocol + Policy):

- **Transcript binding**: sigB covers MessageA and the chosen suite, so modifications to supportedSuites[], keyShares[], or policy fail verification.
- **Negotiation integrity**: Responder MUST select a suite that Initiator offered and for which a key share exists; otherwise reject missingKeyShare.
- **Replay control**: handshakeId is derived from nonces and cached to reject duplicates within a window.
- **Downgrade resistance**: timeout-triggered *default-mode* fallback is disallowed; only whitelisted local capability

errors may fallback under default policy. Strict policy offers either Compliance Mode (no classic establishment) or Bootstrap-Assisted Mode (classic control-only bootstrap with mandatory PQC rekey before business data).

- **Rate limiting:** per-peer fallback cooldown prevents rapid downgrade cycling.
- **Legacy gating:** legacy P-256 acceptance requires an authenticated channel or an existing trust record.

Out-of-Scope (Baseline Claims):

- **Traffic analysis** from size/timing metadata is not assumed in the baseline threat model. SBP2 is an optional mitigation evaluated separately in Section 7, and its overhead/privacy tradeoffs are reported without changing the core security claims.

Trust Assumptions:

- 1) The initial pairing ceremony occurs over a trusted out-of-band channel (e.g., visual PIN comparison on both device screens)
- 2) The device Keychain provides integrity guarantees for stored identity keys
- 3) The Secure Enclave (when available) provides hardware-backed key isolation that resists software-level extraction
- 4) Users can visually verify device identity during pairing (no blind trust)

Out of Scope:

- Side-channel attacks on cryptographic implementations
- Physical attacks on Secure Enclave hardware
- Compromise of the operating system kernel
- Social engineering attacks on users

3.3 Protocol Message Flow

Fig. 1 illustrates the handshake message sequence and transcript coverage. Supplementary Table S10 enumerates suite-negotiation outcomes under the two-attempt strategy.




Fig. 1. Handshake sequence with transcript coverage and policy visibility.

MessageA carries the `supportedSuites[]` (offered suites) and `keyShares[]` arrays (up to two entries to bound size); MessageB selects a suite and binds MessageA via `sigB`. Both sides verify that the selected suite was offered and that a matching key share exists, preventing negotiation mismatch. Under the two-attempt strategy (Section 6.7), PQC and classic are not co-offered in a single MessageA; when fallback is permitted, the Initiator retries with a classic-only MessageA after a whitelist of locally generated capability and negotiation failures. Under strict policy, an optional Bootstrap-Assisted recovery path may transiently establish a classic-only control channel to refresh paired KEM identity keys, but business traffic remains gated until a successful PQC rekey. Explicit key confirmation uses two small Finished frames to avoid half-open state. Full wire-format definitions, field validation rules, and encoding rationale appear in the Supplementary Materials (Wire Format and Validation Rules).

Failure Semantics:

- All failures trigger `HandshakeContext.zeroize()` before error propagation
- Security events are emitted for audit logging
- No partial state is retained after failure

Implementation alignment (artifact snapshot, 2026-01-23).

In the current code revision, connection orchestration for the remote desktop/file-transfer path is route-aware and authentication-gated: local discovery can prefer USB-capable host endpoints, control-channel connections are promoted to “connected” only after handshake authentication, discovery advertisements expose an explicit `hs_soa` capability bit for optional SOA negotiation, and active file-transfer routing prioritizes live authenticated sessions before fallback candidates. These implementation surfaces are enumerated in the Supplementary Artifact Map.

3.4 CryptoProvider Protocol

The CryptoProvider protocol defines a unified interface for all cryptographic operations required by the handshake and session layers [21]. This abstraction enables transparent substitution of underlying implementations without modifying caller code.

The `SigningKeyHandle` abstraction resolves the apparent conflict between “private key as Data” and “Secure Enclave keys never leave hardware.” Software keys are one variant; hardware-backed keys use a reference or callback, ensuring the protocol interface does not assume exportability.

The protocol mandates `Sendable` conformance, ensuring thread-safe usage across Swift’s structured concurrency model. Each provider exposes its tier classification (`nativePQC`, `liboqsPQC`, or `classic`) and active cipher suite, enabling runtime introspection for logging and policy enforcement.

3.5 CryptoSuite Wire Format

A **suite** defines the handshake tuple (`KEM`, `SIG`, `AEAD`, `KDF`) used for the KEM-DEM envelope and transcript binding. Data-plane AEAD is negotiated separately and is fixed to AES-256-GCM in v1 (Supplementary Table S7). Suite identifiers use a compact 16-bit wire ID with reserved ranges for hybrid, PQC, classic, and experimental suites to enable forward compatibility. The full range table, encoding rules, and X-Wing notes appear in the Supplementary Materials (Suite Identifiers and Wire Format). Unknown suite identifiers are parsed as `unknown(wireId)` for audit visibility and forward compatibility; negotiation




Fig. 2. SkyBridge Compass system architecture showing device-local trust boundary (dashed) and policy guard (diamond).

rejects unknown values if they appear in `keyShares` or as `selectedSuite`.

3.6 Provider Factory and Selection

The `CryptoProviderFactory` implements deterministic provider selection based on runtime capability detection and policy configuration:

The selection algorithm proceeds as follows:

- 1) **Capability Detection:** Query the environment for native PQC availability and `liboqs` presence.
- 2) **Policy Application:** Match detected capabilities against the requested policy.
- 3) **Provider Instantiation:** Create the appropriate provider instance for the detected environment and policy.
- 4) **Event Emission:** Emit a `SecurityEvent` recording the selection decision.

For `preferPQC` policy, the precedence order is `ApplePQCProvider` → `OQSPQCProvider` → `ClassicProvider`. This order is deterministic. The `requirePQC` policy returns an `UnavailablePQCProvider`. It throws on all operations if no PQC implementation is available.

Terminology alignment. We use two orthogonal knobs: provider selection (`preferPQC/requirePQC`) and handshake fallback policy (`default/strictPQC`). In our evaluation, `default` corresponds to `preferPQC` with evented classic fallback permitted, while `strictPQC` requires PQC with either (i) Compliance Mode (no classic establishment) or (ii) Bootstrap-Assisted Mode (control-only bootstrap plus mandatory PQC rekey).

3.7 KEM-DEM Authenticated Encryption Format

We use a KEM-DEM envelope ($\text{KEM} \rightarrow \text{HKDF} \rightarrow \text{AEAD}$) as a compatibility layer when native HPKE is unavailable, binding role and transcript into the key schedule with explicit length limits for DoS protection. On Apple 26+ platforms, we use native CryptoKit PQC primitives (ML-KEM/ML-DSA) within the same envelope to preserve a single wire format across platforms. Full header/body formats, versioning, and length limits are provided in the Supplementary Materials (KEM-DEM Envelope).

Why not direct RFC 9180 HPKE everywhere (v1)? Three engineering constraints motivate this compatibility layer in v1: (i) a unified transcript/signature binding surface across mixed provider stacks, (ii) strict backward compatibility with deployed v1 bridge behavior, and (iii) deterministic parser-level DoS checks before provider dispatch. When native HPKE is available (Apple 26+), our v2 path uses RFC 9180 HPKE directly; thus the v1 KEM-DEM envelope is transition code, not a claim of new cryptographic primitive design.

4 HANDSHAKE STATE MACHINE

4.1 Design Principles

The handshake subsystem addresses three critical requirements:

- 1) **Race Condition Prevention:** Concurrent message arrival and timeout expiration must not cause double-resume of continuations.
- 2) **Sensitive Material Protection:** Ephemeral private keys and shared secrets must be zeroized on all exit paths.

- 3) **Observability:** All state transitions and failures must emit structured events.

4.2 Actor-Isolated Architecture

The `HandshakeDriver` actor manages handshake state with compile-time data race prevention:

The `HandshakeContext` actor isolates sensitive cryptographic material:

Reentrancy Considerations. Swift actors permit method re-entry at suspension points (`await`). While actor isolation prevents data races, it does not prevent interleaving of method executions. Our `HandshakeDriver` mitigates reentrancy risks by routing completion through `finishOnce(with:)`, buffering early results, canceling timeouts on completion, and advancing state before `await` points.

Await-window limitation. The `handleMessageB` method awaits `cryptoProvider` operations inside the actor. If a new message arrives during this await, the actor could theoretically process it before the current operation completes. We mitigate this by advancing state to `.processingMessageB` before the await, copying immutable inputs, validating a monotonic epoch counter after resumption, and emitting security events on unexpected transitions.

This approach provides structural protection against reentrancy hazards rather than relying on timing assumptions. Formal model checking of actor await-window interleavings is outside the scope; we mitigate and regression-test this class via epoch checks and targeted fault injection (Table 8; Supplementary Table S6).

4.3 State Transitions

Fig. 3 illustrates the handshake state machine for both roles.

4.4 Simultaneous-Open Arbitration (SOA)

To remove nondeterminism when both peers initiate concurrently, we use an optional `MessageA` extension (SOA TLV) without adding handshake round-trips. The extension carries `soaVersion(1B)`, `initiatorPeerId(32B)`, `targetPeerId(32B)`, `attemptId(16B)`. Arbitration compares the tuple `(initiatorPeerId, attemptId)` with lexicographic *byte* ordering (not string ordering); the smaller tuple wins.

Compatibility is capability-gated (`hs_soa=1` in discovery TXT). If absent, peers stay on the legacy `MessageA` path. The parser keeps `extensionsRaw` bytes (including unknown TLVs) and transcript/hash binding is computed over the original encoded bytes, so decode/re-encode normalization cannot desynchronize signatures. A local outgoing attempt is superseded only after: (i) `MessageA` signature verification succeeds, (ii) `targetPeerId==localPeerId`, (iii) `initiatorPeerId==expectedRemotePeerId`, and (iv) the same canonical `pairKey` is selected. Runtime controls combine `single-flight` (one pending winner per `pairKey`), `established guard` (new attempts return `already_connected`), and `bounded supersede rate` (3 attempts / 60 s). The pending arbitration window is 10 s; `established guard` is released immediately on disconnect.




Fig. 3. Handshake state machines for Initiator (top) and Responder (bottom).

4.5 Double-Resume Prevention

The `finishOnce` method provides a single convergence point for handshake completion, handling the race between timeout expiration and message arrival by: (1) canceling the timeout task on any completion, (2) guarding continuation access with immediate `nil` assignment, and (3) buffering early results for late continuation establishment.

4.6 Secure Zeroization

Swift's standard `Data` and `Array` types use copy-on-write (COW) semantics, which can leave uncontrolled copies of sensitive material in memory. The `SecureBytes` class addresses this by using manual `UnsafeMutableRawPointer` allocation to maintain exclusive ownership, preventing COW-induced secret proliferation. All failure paths explicitly call `context.zeroize()` before error propagation, with `deinit` zeroization as defense-in-depth. On Darwin platforms, zeroization uses platform-supported secure wipe primitives (e.g., `explicit_bzero` or `memset_s`) where available. We avoid exporting secrets as `Data` when possible, but acknowledge that OS-level snapshots, API boundary copies, or privileged inspection can still retain residues; zeroization is a defense-in-depth measure rather than a sole security boundary.

4.7 Transcript Binding and Signature Separation

The protocol uses domain-separated signatures to prevent cross-message and cross-session replay attacks.

Signature A (MessageA) is computed by the Initiator over the `MessageA`-local fields, binding suite preference, policy intent, and identity context before any responder data is received:

For Signature A, the initiator signs `SkyBridge-A`, protocol version, canonical suite/key-share set, client nonce, capabilities, policy, and initiator identity key. Optional Secure Enclave binding uses `SkyBridge-SE-A`. When SOA extensions are present, the signature/transcript includes the raw extension container bytes (unknown TLVs preserved as-is), not a decoded/re-encoded structure.

Signature B (MessageB): The Responder signs a transcript that commits to MessageA:

For Signature B, the responder signs `SkyBridge-B`, transcriptA, selected suite wire ID, responder key share, server nonce, payload hash, and responder identity key. Optional Secure Enclave binding uses `SkyBridge-SE-B`.

Key Derivation: Nonces, transcripts, suite binding, and *directional separation* are included in the KDF, and the salt is transcript-bound (non-empty):

KDF input concatenates the `SkyBridge-KDF` label, suite wire ID, transcript hashes, and nonces. The salt uses the domain string `SkyBridge-KDF-Salt-v1|`. Direction labels separate traffic classes:

- `initiator_to_responder`
- `responder_to_initiator`

This structure is designed to ensure: (1) sigB commits to MessageA, so any tampering breaks verification; (2) domain separators prevent cross-protocol attacks; (3) SE signatures are session-bound and non-replayable.

5 SECURITY MODEL AND GUARANTEES

5.1 Security Model

We assume an active network adversary with full control of the discovery channel (drop, delay, reorder, duplicate, modify, and inject messages) but without breaking standard cryptographic assumptions. The attacker can replay prior handshakes and attempt downgrade by suppressing PQC-related fields. We assume the out-of-band pairing ceremony is trusted and that device key storage (Keychain/Secure Enclave) preserves key integrity. The attacker may compromise a peer after pairing and learn long-term signing key material, but cannot retroactively forge signatures for past sessions. We do not claim forward secrecy for KEM-based suites in v1: ML-KEM encapsulation targets a peer’s long-term KEM public key obtained during pairing. Accordingly, post-quantum confidentiality against store-now-decrypt-later adversaries holds only while long-term KEM private keys remain uncompromised; later compromise of the corresponding KEM private key can enable decryption of recorded sessions.

This v1 choice is intentional for migration: it binds PQC establishment to paired long-term KEM identity material that can be provisioned and audited out-of-band, enabling incremental rollout across heterogeneous fleets without introducing an additional interactive key-registration step. To reduce post-compromise exposure while remaining bridge-compatible, we provide v2 suites that mix the static KEM secret with an ephemeral contribution under transcript binding (Proposition P2) and evaluate the incremental overhead (Supplementary Table S22).

Assumptions:

- The pairing ceremony delivers the correct peer identity key (PIN or visual verification is not bypassed).
- Pinned identity keys stored in Keychain/Secure Enclave retain integrity across normal device operation.

- Devices remain uncompromised (no root or malware control) during handshake execution.

Under concurrent-open threat conditions, the attacker may inject or replay MessageA variants to try to force cancellation of a valid local attempt. The SOA path treats such attempts as unauthenticated control input until signature verification and peer-binding checks complete; therefore cancellation is only admissible for authenticated, correctly bound attempts within the same canonical `pairKey`.

If assumptions fail:

- If pairing is subverted and the pinned identity key is replaced, peer authentication is no longer meaningful; the system must require re-pairing and treat the session as untrusted.
- If device integrity is lost, confidentiality and key secrecy cannot be guaranteed; the system should abort and emit high-severity events such as `identityMismatch` or `handshakeFailed`.

To keep boundary handling operationally explicit, Supplementary Table S38 enumerates boundary-stress cases (pairing social engineering, endpoint compromise classes, and audit-log abuse), their in/out-of-model status, expected event signals, and required operator actions.

5.2 Security Goals, Formal Definitions, and Proof Sketches

Evidence layering. Our mechanized proofs establish security properties for an abstract handshake state machine and its transcript binding. Implementation conformance is supported by property-oriented tests and audit-signal fidelity checks that map each proof obligation to concrete code paths and artifacts (Supplementary Table S40). Full code-level verification of the Swift implementation is outside the scope of this paper.

Definitions.

- **D1 (Forced downgrade / forced fallback).** A downgrade occurs if a business-capable session is established with a classic suite when strict policy requires PQC, or if classic fallback is accepted under the default policy without satisfying the local-only fallback gate (Definition D2) and without emitting the corresponding audit event (Property G8).
- **D2 (Local-only fallback gate, default policy).** Fallback is permitted only as an *initiator-side attempt switch* when (i) the session targets an already paired peer (pinned trust record) and (ii) the triggering failure is a locally generated capability/instantiation unavailability detected *pre-send* at the on-wire send boundary (i.e., before emitting the first handshake flight / MessageA in the current attempt); timeout and any network-derived parse/verification failure are fallback-ineligible by definition.
- **D3 (Audit evidence).** The minimal sufficient evidence set is the tuple {event type, timestamp, deviceId, suiteId, failure reason code, policy mode, fallback flag, cooldown remaining} (Property G8).
- **D4 (Bootstrap-Assisted strict recovery gate).** In strict policy, a temporary classic control channel is allowed only for paired KEM-key recovery triggers (`missingKEM`, `staleKEMRecovery`); non-control business messages are blocked until PQC rekey succeeds.

Lemma L1 (Timeout-triggered fallback impossibility). Consider any migration policy that permits classic fallback on timeout or other unauthenticated network-derived failures. Under the active attacker model in Section 5 (drop/delay/reorder), the adversary can force repeated timeout observations on the PQC path and thereby force classic establishment. *Proof sketch:* the attacker suppresses or delays PQC-path messages past the timeout boundary; if timeout is fallback-eligible, each attempt has a reachable classic edge under attacker control.

Corollary C1 (Necessary fallback gate condition). To retain compatibility fallback without attacker-forced downgrade, fallback eligibility must exclude timeout/network-derived failures and be restricted to attacker-non-inducible local causes (Definition D2), or classic fallback must be disabled entirely.

We state the core properties as theorem/proposition-level claims tied to standard assumptions and protocol mechanics. To reduce vacuity risk, the released Tamarin artifact models explicit classic/PQC/v2 traces, explicit fallback-gate and downgrade events, and includes existence lemmas in addition to safety lemmas (formal/skybridge_v2.spthy) [32]. This revision additionally verifies timeout-fallback denial, cooldown rate limiting, downgrade-context evidence, and bootstrap control-only rekey ordering; their implementation-level traceability anchors are summarized in Supplementary Table S40.

Formal security definitions (experiment style).

- Exp_{auth} : adversary wins if it causes a paired endpoint to accept an unauthenticated peer identity.
- Exp_{neg} : adversary wins if `selectedSuite` is accepted outside the signed/advertised offer context.
- Exp_{dg} : adversary wins if classic establishment occurs in violation of strict policy, or under default policy without local fallback gate + downgrade event.
- $\text{Exp}_{\text{fs-hybrid}}$: for v2, adversary wins if it computes session keys without both static KEM contribution and ephemeral contribution under transcript binding.

Assumption set.

- KEM IND-CCA security for negotiated KEM suite (ML-KEM/X-Wing profile).
- Signature EUF-CMA security for protocol signatures (ML-DSA-65 / Ed25519 as configured).
- HKDF-SHA256 as a PRF in the used extraction/expansion domain.

Theorem T1 (Authentication). Under EUF-CMA and trusted pairing pins, advantage in Exp_{auth} is negligible up to bounded implementation error terms (e.g., policy misconfiguration outside the model).

Theorem T2 (Negotiation integrity). Under EUF-CMA and transcript binding, advantage in Exp_{neg} is negligible.

Proposition P1 (Policy-gated downgrade). Under the threat model in Section 5, timeout/drop/reorder adversaries cannot force successful classic establishment in strict policy, and cannot produce silent classic fallback in default policy without the local gate event.

Proposition P2 (v2 hybrid FS semantics). For v2 sessions, compromise of long-term KEM private key alone is insufficient to reconstruct the derived session root unless the ephemeral contribution secret is also recovered. This is an intermediate secrecy guarantee and is intentionally weaker than full ephemeral-PFS.

Property G1 (Negotiation Integrity): The selected suite must be among the Initiator’s offered suites and have a corresponding key share.

Proof sketch:

- `MessageB` includes `selectedSuite`, and `sigB` covers `MessageA`.
- Any modification to the offered suite list or key-share list breaks `sigB` verification.
- The Responder validates `selectedSuite` against the offered suites and checks that a matching key share is present.

Property G2 (Mutual Authentication for Paired Peers): For previously paired devices, the peer identity is authenticated and bound to the handshake transcript.

Proof sketch:

- `identityPubKey` is pinned during OOB pairing.
- Signatures in `MessageA` and `MessageB` over transcript data must verify under the pinned key; otherwise `identityMismatch` aborts the handshake.

Property G3 (Session Key Secrecy): The derived session keys remain secret under standard KEM/DH and signature security assumptions. For KEM-based PQC suites in v1, we do not claim post-compromise forward secrecy: confidentiality holds only while long-term KEM private keys remain uncompromised. In key-control terms, v1 KEM establishment is initiator-contributory to a responder static key; v2 adds explicit responder ephemeral contribution to reduce unilateral key-control concentration.

Proof sketch:

- The shared secret is derived from KEM decapsulation or DH, and keys are extracted via HKDF with transcript-bound salt.
- An attacker without the private key material cannot compute the shared secret or derive session keys.

Property G4 (Replay Resistance): Replayed handshakes do not establish a session.

Proof sketch: `handshakeId` derives from fresh nonces and suite identifiers; replay detection caches recent IDs and rejects duplicates within the window.

Property G5 (Strict-Policy Consistency): Strict policy never permits a business-capable classic session. Compliance Mode has no classic path. Bootstrap-Assisted Mode may use a transient classic control channel, but business traffic is blocked until PQC rekey completes.

Proof sketch:

- `TwoAttemptHandshakeManager` enforces a policy gate before any default-mode fallback attempt.
- Strict compliance removes classic establishment edges.
- Strict recovery trigger-limits bootstrap, emits explicit `handshake_fallback` events, and keeps messaging control-only until PQC rekey succeeds.

We validate this behavior using policy downgrade benchmarks (Fig. 4) and bootstrap control-gating tests.

Property G6 (Safe Fallback under default policy): Fallback can only occur for already paired peers and only for a whitelisted set of locally generated *pre-send* capability/instantiation errors; timeout and network-derived failures never trigger default-mode

fallback.

Proof sketch: Fallback is gated on a pinned trust record and a pre-send whitelist evaluated at the initiator-side attempt boundary (before emitting MessageA); whitelisted causes are restricted to stable local capability predicates (Supplementary Table S34). `timeout` is explicitly blocked and per-peer cooldown limits repeated downgrades. *Proposition (No network-induced suiteNegotiation fallback):* Under Definition D2, `suiteNegotiationFailed` can trigger default-mode fallback only in the local pre-send branch (no local keyshare constructible for the offered PQC suite). If the failure is observed after remote input, it is treated as network-derived and cannot satisfy the fallback gate. *Proposition (No forced fallback by packet loss):* Under the attacker model in Section 5, an active network attacker who can drop/delay/reorder messages cannot induce a successful classic session by causing timeouts, because timeout-triggered fallback is disallowed and the default-policy fallback gate (Definition D2) requires a locally generated pre-send unavailability for an already paired peer.

Property G7 (Legacy Acceptance Preconditions): Legacy P-256 signatures are accepted only under authenticated pairing or an existing trust record.

Proof sketch: Legacy acceptance is gated by `LegacyTrustPrecondition`; pure network stranger connections fail. This is validated by the precondition test matrix (Table 7).

Property G8 (Auditability): Cryptographic downgrades and exceptional states are observable via a minimal sufficient information set.

Proof sketch: The handshake emits `cryptoDowngrade` and related events with reason, `deviceId`, and cooldown context, and records them into a per-session hash-chain (`AuditTrail`) anchored to the handshake transcript hash, enabling post-hoc verification of policy adherence and event-order integrity.

Audit field requirements:

- **MUST log** (minimal sufficient set): event type, timestamp, `deviceId`, `suiteId`, failure reason code, policy mode, fallback triggered (bool), cooldown remaining.
- **MUST NOT log** (privacy preservation): session keys, nonces, raw signatures, message payloads, user identifiers beyond `deviceId`.
- **Tamper-evidence:** In addition to synchronous event emission, each session maintains a transcript-anchored audit hash-chain (`AuditTrail`). This provides cryptographic tamper evidence for in-session event order/content under non-compromised runtime assumptions. Long-term retention and remote attestation depend on the operating environment and are out of scope; privileged adversaries can truncate host logs.

5.3 Complexity and Overhead Bounds

Let S and K denote implementation bounds for offered suites and key shares (`maxSupportedSuites`, `maxKeyShareCount`); currently $S \leq 8$ and $K \leq 2$. Encoding/decoding work for MessageA/MessageB is:

$$T_{\text{codec}} = O(S + K + |M_A| + |M_B|),$$

with constant-time checks on fixed-length fields and linear passes over suite/key-share vectors.

Handshake cryptographic operation counts are:

- **v1 PQC path:** $1 \times \text{KEM encapsulate} + 1 \times \text{KEM decapsulate} + 2 \text{ signatures} + 2 \text{ signature verifies} + \text{HKDF key schedule}$.
- **v2 FS path:** v1 counts plus one ephemeral contribution generation per side and one additional shared-secret composition step (HKDF extract-expand on static+ephemeral input).
- **Classic path:** one DH-style encapsulation/open path (KEM-DEM abstraction), signatures/verifications, and HKDF key schedule.

Closed-form byte overhead between v1 and v2 is dominated by the extra contribution fields. Let c_I denote the initiator contribution carried in MessageA and $c_R^{(v1)}, c_R^{(v2)}$ denote responder-share fields in MessageB for v1 and v2:

$$\Delta|M_A| = 2 + |c_I|, \quad \Delta|M_B| = |c_R^{(v2)}| - |c_R^{(v1)}|.$$

In current implementation (x25519 contribution, 32 bytes): $\Delta|M_A| = 34$ bytes and $\Delta|M_B| = 32$ bytes.

6 IMPLEMENTATION

Implementation references (source files, line numbers) are consolidated in the Supplementary Materials (Artifact Map).

6.1 Platform Support Matrix

SkyBridge Compass uses a tiered provider model: a native PQC provider when available (e.g., CryptoKit ML-KEM/ML-DSA), a portable liboqs provider, and a classic provider. Provider selection is deterministic and prefers the strongest available tier; Secure Enclave proof-of-possession is optional and orthogonal to the negotiated suite. The full Apple provider matrix and the cross-platform runtime suite-policy matrix are in the Supplementary Materials (Tables S31 and S33).

6.2 Native PQC Integration

The native provider wraps CryptoKit ML-KEM/ML-DSA APIs [16], [23] when the SDK exposes them, enabling platform-backed PQC suites and HPKE-based encapsulation. We treat this as a deployment instance rather than a protocol dependency; integration details are provided in the Supplementary Materials (Native PQC Integration). Specifically, the Apple provider path does not redefine the protocol contract, does not alter the primary evidence base of the paper, and is evaluated only as an implementation capability mapping under existing negotiation and policy semantics.

6.3 Conditional Compilation Strategy

PQC symbols are gated behind build-time flags to keep the codebase buildable on SDKs that lack PQC types and to avoid compile-time failures. The flagging strategy and build settings are detailed in the Supplementary Materials (Conditional Compilation).

6.4 Security Event Emission

All cryptographic decisions emit structured audit events with backpressure and rate limiting; implementation details are in the Supplementary Materials (Security Event Emission).

TABLE 3
Security Contract (Properties, Enforcement, Evidence).

Property	Enforced at	Evidence
G1 Negotiation integrity	Suite/keyshare validation + transcript-bound sigB	Fig. 1, HandshakeDriverTests
G2 Mutual authentication	Identity pinning + signature verification	Supplementary Table S6 (wrong signature), Table 7
G3 Session key secrecy	KEM/DH + HKDF with transcript-bound salt	Property 1, Supplementary Materials (ML-DSA Key Sizes)
G4 Replay resistance	handshakeId cache + nonce binding	Property-Oriented Testing (Section 7)
G5 Strict-policy consistency	TwoAttempt policy gate + bootstrap control-only gate + mandatory rekey to PQC	Fig. 4, PolicyDowngradeBenchTests, BootstrapAssuranceTests
G6 Default safe fallback	Whitelist/blacklist + local capability gate + cooldown	Fig. 5, Supplementary Table S6
G7 Legacy acceptance precondition	LegacyTrustPrecondition	Table 7
G8 Auditability	Security event emission + transcript-anchored AuditTrail hash-chain	Fig. 6, Table 8

6.5 Secure Enclave Integration

For hardware-backed signing, a `SigningCallback` protocol allows Secure Enclave keys without exposing private material. Availability, fallback behavior, and key-type details are provided in the Supplementary Materials (Secure Enclave Integration).

6.6 Signing Key Hierarchy

The system uses two complementary signatures: (1) protocol signatures via the negotiated suite (Ed25519 for classic, ML-DSA-65 for PQC) to bind peer identity to the transcript, and (2) optional Secure Enclave P-256 signatures for device proof-of-possession. Protocol signatures are mandatory; Secure Enclave PoP is optional and provides a hardware-backed signal. Key hierarchy details and verification rules are provided in the Supplementary Materials (Signing Key Hierarchy).

6.7 Pre-Negotiation Signature Algorithm Selection and Two-Attempt Strategy

A fundamental challenge is that `sigA` must be generated before negotiation completes, yet the signature algorithm should match the negotiated suite. We resolve this with pre-negotiation signature selection and a two-attempt strategy. Each attempt is homogeneous (PQC-only or classic-only), and the signature algorithm is chosen based on the offered suites. The initiator first attempts PQC, then falls back to classic only for already paired peers and only for a whitelist of benign, locally generated *pre-send* capability/instantiation errors; timeout, signature verification, identity mismatch, replay, and other network-derived failures never trigger fallback. This policy is enforced by the handshake manager and logged as security events. Detailed error lists and type-level enforcement are provided in the Supplementary Materials (Two-Attempt Strategy). For strict policy deployments, we expose two operator-selectable modes: *Compliance Mode* (hard fail when PQC cannot be established) and *Bootstrap-Assisted Mode* (allow one temporary classic control channel only for paired KEM-key recovery, then require PQC rekey before business traffic).

Compliance Mode is fail-closed and is the reference point for our strict downgrade-resistance claims. Default Mode prioritizes availability and therefore allows a bounded, policy-audited classic fallback for already paired peers under the local-only gate (D2); persistent PQC blocking is treated as an availability attack rather than a silent downgrade.

Fallback Boundary Conditions (Explicit Scope):

TABLE 4
Default-policy fallback classes and remote-inducibility boundary.

Reason class	Default-policy outcome
Local capability/instantiation unavailability (paired trust record, pre-send; e.g., local <code>suiteNegotiationFailed</code> due to missing paired KEM keyshare)	ALLOW classic fallback + mandatory <code>cryptoDowngrade</code> event
Network-derived / after-send failures (<code>timeout</code> , remote negotiation mismatch, signature/identity/replay failures)	DENY fallback; emit failure event only

- *We defend against:* (1) Attacker-induced timeout forcing downgrade (blocked by timeout blacklist); (2) Rapid downgrade cycling via repeated failures (per-peer cooldown; default 5 minutes); (3) Signature/identity attacks masquerading as capability failures (verification failures are blacklisted).
- *We do not defend against:* (1) Persistent PQC blocking under Default Mode (availability vs. security trade-off; Compliance Mode is fail-closed); (2) Legitimate PQC provider failures causing temporary classic fallback under Default Mode.
- *Cooldown rationale:* We set a conservative default of 5 minutes to bound downgrade cycling while limiting availability impact; the cooldown is a tunable operator knob (ablation in Supplementary Table S5). Operators requiring zero classic establishment can select strict Compliance Mode.

Attempt State and Identifiability:

- *Shared across attempts:* pinned identity fingerprint and pinned KEM public keys (trust record), and per-peer fallback cooldown state.
- *Not shared across attempts:* nonces, transcript hashes, handshakeId/replay tags, ephemeral key shares, and session keys; each attempt runs a fresh state machine.
- *Identifiability:* a passive observer can distinguish a classic fallback attempt from a PQC attempt by wire format. We do not transmit an explicit attempt counter; strict Compliance Mode removes classic attempts entirely when capability privacy outweighs compatibility.

Security Events:

Under default policy, classic fallback emits event “cryptoDowngrade” with reason, deviceId, and cooldown fields. Strict

Bootstrap-Assisted recovery emits event “handshake_fallback” (bootstrap-assisted mode) with explicit trigger/result fields. This distinguishes strict compliance (fail-closed) traces from strict recovery traces with mandatory PQC rekey.

7 EVALUATION

7.1 Experimental Setup

We evaluate SkyBridge Compass along two primary tracks: **security-centric evidence** (fault injection, downgrade suppression, legacy precondition enforcement, and auditability) and **cost-centric evidence** (handshake latency, wire overhead, provider selection overhead, and data-plane throughput).

Environment: All experiments are run on Apple Silicon (M1/M3 class) with macOS 26.x, where CryptoKit PQC is available when the SDK exposes ML-KEM/ML-DSA types [16], [23]. Build configuration uses Release (-O) with non-essential logging disabled. Timing uses a monotonic clock (ContinuousClock) with 10 warmup runs. The *core gate* (Classic/liboqs/liboqs-v2-FS) uses $N=1000$ per configuration. Apple PQC/X-Wing are collected in a separate *contrast* pass ($N=5000$) and are reported as non-gating diagnostics. Reported percentiles (p50/p95/p99) are computed directly from samples; binary outcomes are reported with Wilson 95% confidence intervals.

Repeatability: We support multi-batch repeatability by restarting the benchmark test process for each batch. Supplementary Tables S25–S26 report the number of observed batches (B) and (when $B \geq 2$) mean and 95% confidence intervals across batches. To reproduce multi-batch CIs, re-run the harness with `SKYBRIDGE_BENCH_BATCHES=3–5`.

Reproducibility: We ship an opt-in benchmark suite under `Tests/SkyBridgeCoreTests/` that emits date-stamped CSV artifacts under `Artifacts/` (pinned by `ARTIFACT_DATE=2026-01-23`) and regenerates all evaluation tables/figures from scripts. For a one-shot run, use `Scripts/run_paper_eval.sh`. The runner supports profile-scoped collection via `SKYBRIDGE_BENCH_PROFILE` and `SKYBRIDGE_BENCH_OUTPUT_SET`; contrast outputs are written to dedicated contrast CSV/JSON artifacts. The complete command reference and CSV artifact specifications are provided in the Supplementary Materials (Reproducibility).

iOS Reproducibility: We ship an on-device micro-benchmark in the iOS app (`Settings → PQC → Self-test/Bench`) that exports schema-v3 JSON artifacts. For the pinned run (`ARTIFACT_DATE=2026-01-23`), each suite uses `warmup=10`, $N=1000$, and three batches ($B=3$). We aggregate exported JSON with `aggregate_ios_microbench.py` into the date-stamped iOS micro-benchmark CSV, then auto-generate the main table and Supplementary device metadata. In this frozen run (iOS/iPadOS 26.2.1), the iOS compatibility-matrix gate passed on both tested device classes (iPad + iPhone), with no schema or negotiation incompatibility observed in the benchmark pipeline.

iOS microbench results: Supplementary Table S21 reports *primitive-level* iOS costs (encap/sign/verify/seal+open), not end-to-end handshake completion; it is therefore not interpreted as a direct substitute for macOS loopback handshake latency. Supplementary Table S20 records per-device metadata for the pinned run.

7.2 Claim Scope and External Validity

We separate *mechanism-level* from *deployment-level* claims. Mechanism-level claims cover negotiation integrity, policy-gated downgrade behavior, deterministic failure semantics, and audit-event completeness under the stated threat model. These claims are tied to the protocol contract and test/formal artifacts, not to Apple-specific APIs. Deployment-level claims cover absolute latency, wire overhead, and provider/runtime effects. They are evaluated only for the reported Apple-platform instance and should not be extrapolated to other OS stacks without remeasurement. To reduce portability ambiguity, we ship:

- a cross-platform contract-consistency JSON artifact (Supplementary Table S32);
- a dedicated external Linux direct-path run (Supplementary Table S18).

Accordingly, we use the real-network results as external-validity evidence rather than ecosystem-wide performance characterization: the current snapshot includes one STUN-coupled non-overlay direct path, overlay-mediated paths, and a dedicated external Linux direct-path run ($n=50$ per payload), and is used to validate directional robustness and failure semantics under deployment constraints.

Evidence overview (main claims):

- Negotiation integrity and transcript binding: formal proof obligations + protocol/property tests (Section 5; `formal/skybridge_v2.spthy`).
- Policy-gated downgrade correctness: decision-matrix validation + forced-error policy bench + migration-threat harness contrast (Fig. 4; Fig. 5; Table S9; Table 5).
- Deterministic failure semantics and event typing: fault injection + spec-to-event trace mapping (Supplementary Table S6; Table 8).
- Performance and overhead characterization: distributional microbench + system-level amortization (Table 9; Fig. 10; Section 7.6).
- External-validity guardrail (deployment instance): controlled simulation + kernel cross-check + fleet dynamics + real-network micro-study (Section 7.7; Supplementary Tables S19, S18, S23).

Supplementary Table S3 provides the full claim-to-evidence map with artifact anchors.

Threats to Validity (Reviewer-facing).

- *Internal:* timing can be affected by runtime noise (scheduler/cache/thermal), so we use Release builds, warmup, distributional reporting, and multi-batch restarts.
- *Construct:* loopback *time-to-ready* and microbench costs are not user QoE; we report `T_connect`, `T_first_frame`, and `T_total` separately.
- *External:* absolute costs are Apple-platform instance results; cross-stack evidence is framed as contract-consistency plus targeted Linux direct-path checks.
- *Conclusion:* audit-signal fidelity is deterministic spec-to-event conformance (not statistical); artifacts are published for independent reruns.

7.3 Security Validation

We validate security properties through fault injection, property-based testing, and audit-signal fidelity analysis. This section

presents the key results; detailed test matrices and methodology are provided in the Supplementary Materials (Regression Testing).

Migration-threat harness (baseline contrast). To evaluate migration behavior under active suppression, we run a targeted harness with four metrics: **FDR** (forced downgrade rate), **SDR** (silent downgrade rate), **ECR** (evidence completeness rate for emitted downgrade events), and **PCR** (policy-decision replayability rate from recorded traces). We compare a naive timeout-fallback baseline against SkyBridge policies under a timeout-attack scenario (N=1000). This baseline models a common availability-first glue policy where timeouts are treated as fallback triggers; Lemma L1 applies to any migration strategy that makes timeout/network-derived failures fallback-eligible.

This contrast provides implementation-level evidence for Lemma L1 and Corollary C1: timeout-eligible fallback is attacker-forcible, while a local-only gate removes that adversarial edge.

Provider/back-end consistency check (non-Apple path). To validate that migration semantics are not tied to Apple-native providers, we rerun the same harness using explicit liboqs and classic provider selection. Supplementary Table S4 reports the provider-matrix results (N=1000 per row).

Fleet-level migration dynamics (heterogeneous capability mix). To quantify migration behavior beyond per-handshake microbenchmarks, we run a controlled fleet simulation with heterogeneous peer capability ratios, per-peer cooldown on/off ablation, and active timeout suppression (N=10,000 sessions per row; peer pool=2,000; timeout attack rate=5%). Supplementary Table S5 reports availability, downgrade exposure, and policy-decision distribution from `Artifacts/migration_fleet_behavior_2026-01-23.csv`.

The fleet-level mean connect metric is an in-process policy/attempt completion cost (not external network RTT) and is reported only to compare policy modes under identical harness conditions.

Failure-Mode Robustness. The actor-isolated handshake driver targets deterministic failure semantics; in our fault-injection tests we observed no unhandled errors, double-resume, or sensitive-material leaks. We inject faults including timeout, malformed framing, invalid signatures, out-of-order delivery, and concurrent cancellation (n=1000 per scenario per policy).

Unknown-suite handling: We parse unknown suite identifiers as an explicit `unknown(wireId)` value and reject them if they appear in `keyShares` or as `selectedSuite`. This prevents silent “unknown → classic” coercions that would otherwise weaken downgrade guarantees.

Supplementary Table S6 reports the full per-scenario robustness and observability matrix and event counts (n=1000 per scenario; strictPQC counts recorded separately in `fault_injection_2026-01-23.csv`).

The delay-within-timeout and delay-exceed-timeout scenarios approximate RTT jitter and loss-induced stalls, while drop and duplicate cases approximate packet loss and reordering. The per-scenario event distributions are recorded in `Artifacts/fault_injection_2026-01-23.csv`.

Downgrade events are quantified separately in the policy bench (Fig. 4) to avoid conflating fallback behavior with transport corruption scenarios.

SOA Arbitration and Compatibility Stress. We add a focused SOA regression bench (`SOAInteroperabilityBenchTests`, N=100 per scenario) covering simultaneous-open convergence, forged-target rejection, established-guard




Fig. 4. Policy guard enforces strictPQC Compliance Mode: fallback events per 1000 runs, macOS 26.x, N=1000 per policy. Error bars show 95% Wilson binomial CI.




Fig. 5. Downgrade decision matrix (policy × error) with explicit allow/deny semantics, macOS 26.x. Each cell is labeled ALLOW or DENY for grayscale readability.

release, supersede rate limiting, and legacy/new MessageA interoperability. Table 6 summarizes the results; raw rows are stored in `Artifacts/soa_interop_2026-02-22.csv`.

The plot in Fig. 4 uses the policy-downgrade dataset. File: `policy_downgrade_2026-01-23.csv`. It demonstrates that strictPQC Compliance Mode never emits fallback events even under forced PQC-unavailable errors; default policy shows non-zero downgrade events.

Fig. 5 encodes the explicit downgrade whitelist/blacklist: only PQC-unavailability and local pre-send suite-selection errors may fallback under default policy; strictPQC Compliance Mode denies classic fallback edges (Bootstrap-Assisted strict recovery is a separate control-only path with mandatory PQC rekey).

Fig. 6 combines fault-injection counts with the downgrade-acceptance event from the policy bench to visualize observability of failures and downgrades across policy modes.

Property-Oriented Testing. Correctness properties are validated with property-based tests using a reproducible seed (0xDEADBEEF):

- KEM-DEM round-trip correctness across all providers (100 trials)
- Signature verification integrity (100 payloads per provider)
- Secure zeroization on deallocation
- State-machine single-resume guarantee under concurrent stress (500 iterations)

A standalone Python decoder (`Scripts/wire_format_sanity.py`) provides non-Swift interoperability validation of wire formats.

TABLE 5
Migration-threat harness under active timeout suppression (N=1000 per row). Data source: Artifacts/
migration_threat_harness_2026-01-23.csv produced by MigrationThreatHarnessBenchTests.

Configuration	Scenario	FDR	SDR	ECR	PCR	Interpretation
SkyBridge default	Timeout attack	0.00	0.00	1.00	1.00	Timeout does not open classic fallback edge
SkyBridge strictPQC	Timeout attack	0.00	0.00	1.00	1.00	Compliance mode remains fail-closed under suppression
Naive timeout-fallback baseline	Timeout attack	1.00	1.00	0.00	0.00	Attacker can force classic with no replayable audit evidence
SkyBridge default	Local unavailability	1.00	0.00	1.00	1.00	Local-only fallback path is explicit and fully evidenced

TABLE 6
SOA arbitration and compatibility stress results (ARTIFACT_DATE=2026-02-22, N=100 per row).

Scenario	Check focus	N	Pass	Pass rate
Simultaneous open single-flight	Deterministic winner via byte-ordered (initiatorPeerId,attemptId)	100	100	1.00
Forged target binding guard	Reject unauthenticated cancellation trigger (targetPeerId mismatch)	100	100	1.00
Established guard release	already_connected while established; immediate reconnect after release	100	100	1.00
Supersede rate limit	Bound supersede to 3/60s; 4th attempt rejected	100	100	1.00
Legacy MessageA compatibility	Missing SOA extension remains valid legacy path	100	100	1.00
Extension passthrough + transcript binding	Unknown TLVs preserved in extensionsRaw; SOA-only wire delta = 91 B, no extra RTT	100	100	1.00




Fig. 6. Failure-mode histogram of security events, macOS 26.x, N=1000 per scenario.

ML-DSA Key Lifecycle. ML-DSA-65 key generation, storage, and sign/verify operations are validated against FIPS 204 [24]. Size-compliance results are reported in the Supplementary Materials (ML-DSA Key Sizes). All property tests (round-trip, tampering detection, key independence) achieve 100% pass rate. Keys are stored in Keychain with device-local accessibility and synchronization disabled (kSecAttrSynchronizable=false).

Legacy Fallback Security. Legacy P-256 signatures are only accepted when a security precondition is satisfied: (1) authenticated out-of-band channel (QR code, PAKE/PIN), or (2) existing TrustRecord with legacyP256PublicKey. Pure network stranger connections are rejected to prevent downgrade attacks.

Property G7 (Legacy Fallback Security Precondition): A legacy P-256 signature SHALL only be accepted when a security precondition is satisfied. Pure network stranger connections SHALL be rejected. Validated with 12 test cases covering all precondition combinations (100% coverage of the precondition state space).

Audit-Signal Fidelity. Security events provide deterministic, protocol-level audit signals of failures and policy decisions. Under non-compromised runtime assumptions, the per-session AuditTrail hash-chain provides tamper-evident ordering; long-




Fig. 7. Audit-signal event trace (representative timeout case) showing attacker action, policy gate, and emitted event.

term retention depends on the operating environment. Fig. 7 illustrates a representative timeout trace showing attacker action, policy gate response, and emitted event.

Evaluation approach: We evaluate audit-signal fidelity using a spec-to-event mapping consistency metric, *not* a statistical classifier. Each fault-injection scenario deterministically maps to exactly one expected trace template (event sequence), e.g., timeout → handshakeFailed=1, cryptoDowngrade=0. Match rate (MR)=1.00 means all N iterations matched the expected template; spurious rate (SR)=0.00 means no unexpected events were emitted. This is regression consistency against the protocol specification, not ML classification performance.

Table 8 summarizes results across all fault-injection scenarios: all scenario classes achieve 1.00 match and 0.00 spurious rates, confirming correct event semantics.

Due to space constraints, the complete regression test matrix, TLV encoding validation, and implementation component details are provided in the Supplementary Materials (Regression Testing and Artifact Map).

7.4 Resource Amplification and Throttling

We separate adversarial cost vectors into (i) pre-authentication parser cost and (ii) post-authentication retry/cooldown pressure. Pre-auth cost is bounded by strict field-length checks, overflow-

TABLE 7

Legacy Fallback Precondition Test Results. All tests executed via LegacyFallbackPreconditionTests. Tests validate Property G7: Legacy Fallback Security Precondition.

Scenario	Precondition	Expected Result	Actual Result
Pure network stranger	None	Reject	OK Reject
QR code pairing (verified)	authenticatedChannel	Allow	OK Allow
PAKE pairing (verified)	authenticatedChannel	Allow	OK Allow
Existing TrustRecord with legacy key	existingTrustRecord	Allow	OK Allow
Existing TrustRecord without legacy key	None	Reject	OK Reject
Network discovery	None	Reject	OK Reject
Unverified authenticated channel	None	Reject	OK Reject

TABLE 8

Audit-Signal Fidelity Summary. Match/spurious rates reflect *deterministic* protocol event mapping (each scenario maps to exactly one expected trace template), not statistical classification. Perfect scores indicate regression consistency against the protocol specification.

Scenario Class	Expected Signal	N	Match (MR)	Spurious (SR)
Drop/Timeout (default)	handshakeFailed=1, cryptoDowngrade=0	3000	1.00	0.00
Drop/Timeout (strictPQC)	handshakeFailed=1, cryptoDowngrade=0	3000	1.00	0.00
Corrupt/Wrong Sig	handshakeFailed=1, cryptoDowngrade=0	6000	1.00	0.00
Ordering Benign	handshakeFailed=0, cryptoDowngrade=0	6000	1.00	0.00
PQC unavailable (default)	cryptoDowngrade=1	1000	1.00	0.00
PQC unavailable (strictPQC)	cryptoDowngrade=0	1000	1.00	0.00

safe parsing, and bounded key-share cardinality before cryptographic operations are invoked. Post-auth pressure is bounded by per-peer fallback cooldown and rate-limited security event emission.

These controls are intended to bound resource amplification from malformed inputs and retry storms without introducing silent downgrade paths. We do not claim Internet-scale saturation resistance from this revision alone; adversarial burst-throughput characterization remains future work.

7.5 Performance Benchmarks

Table 9 consolidates *core gate* performance metrics. We measure end-to-end handshake completion time using an in-memory loopback transport to isolate cryptographic and state-machine overhead. Each core configuration reports distributional statistics after 10 warmup runs (N=1000 per configuration). Apple PQC/X-Wing are reported separately as non-gating contrast data.

v1/v2 Bridge Overhead. To isolate the protocol-v2 cost, we run a controlled liboqs v1/v2 bridge comparison with identical hardware/runtime settings (N=1000, warmup=10). Supplementary Table S22 reports mean/p95 latency, RTT, and payload-handshake bytes. Consistent with the wire-format analysis in Section 5.3, v2 adds a fixed +75 B payload-handshake overhead (MessageA +43 B, MessageB +32 B), while latency impact remains in the low-millisecond range.

Latency Distribution. Fig. 8 shows latency percentiles across configurations. Classic handshakes remain near 2.232 ms (2.503 ms at p99). The liboqs v2 FS configuration remains around 4.572 ms at p99. Extended statistics (mean/p50/p95/p99, std) are available in Supplementary Tables S12–S13 (reported in the supplementary PDF).

Wire-Format Breakdown. The handshake payload size increases from Classic (687 B) to PQC (12002 B for liboqs; 12077 B for liboqs v2 FS in the core gate), dominated by ML-DSA




Fig. 8. Handshake latency percentiles for the core gate configurations (Classic, liboqs PQC, liboqs PQC v2 FS), warmup=10 and N=1000 per configuration. Data source: `handshake_bench_2026-01-23.csv`. Values align with Table 9 and Supplementary Table S12.

signatures (3309 B) and ML-KEM ciphertexts (1088 B). Apple PQC/X-Wing payloads are reported in supplementary contrast artifacts. Fig. 9 shows the per-field breakdown. **Path-MTU implication.** The payload-only core-gate PQC handshake (12002 B for liboqs; 12077 B for liboqs v2 FS) spans multiple transport segments under common MTUs. In our design, handshake messages run over stream transports (TCP/QUIC stream), so segmentation/reassembly is delegated to transport rather than custom handshake-layer fragmentation; the practical risk is latency inflation on lossy paths, which is why we report controlled impairment, kernel-level, and real-network evidence in Section 7.7.

Note: In the 2026-01-23 snapshot, X-Wing hybrid KEM adds 128 B to MessageA’s keyshare (1216 B vs 1088 B for ML-KEM-768). Measured handshake sizes are reported in Supplementary Table S14 (MessageA.XWing/MessageB.XWing).

TABLE 9

Performance Summary. All benchmarks on Apple Silicon (M1/M3), macOS 26.x, N=1000 iterations after 10 warmup runs. Wire Size counts payload-only handshake bytes (MessageA + MessageB + 2×Finished); loopback wire sizes including transport overhead are reported separately in Table 2 and Supplementary Table S24. Data-plane AEAD is fixed to AES-256-GCM in v1, so throughput is independent of the negotiated handshake suite; throughput measured post-handshake on 1 MiB payloads.

Configuration	Handshake Latency		RTT		Wire Size (bytes)	Throughput (GB/s)
	mean (ms)	p95 (ms)	p50 (ms)	p95 (ms)		
Classic	2.23	2.42	0.47	0.53	687	3.7
liboqs PQC	2.74	3.40	0.82	1.37	12,002	3.7
liboqs v2 FS	3.48	4.18	1.15	1.70	12,077	3.7




Fig. 9. Handshake message size breakdown (signature, keyshare, identity fields, framing overhead) for MessageA/MessageB across Classic and PQC configurations. Counts reflect MessageA/B payloads only. The *payload-only* handshake size adds two 38 B Finished frames (accounted for in Table 9); detailed field sizes are reported in Supplementary Table S14. Data source: `message_sizes_2026-01-23.csv`.

Data-Plane and Provider Selection. Post-handshake symmetric AEAD (AES-256-GCM) throughput reaches 3.7 GB/s for 1 MiB payloads on Apple Silicon, with CPU cost of 0.25 ns/byte. In v1 the data-plane AEAD is fixed to AES-256-GCM and recorded separately from the handshake suite; thus throughput is independent of the negotiated KEM/signature suite. Provider selection overhead is sub-microsecond for cached paths (0.17 μ s p50) and under 6 μ s even for fallback scenarios. Detailed breakdowns are provided in the artifact CSV files.

7.6 System-level Impact and Amortization

The preceding microbenchmarks isolate handshake and cryptographic costs. To align evaluation with the *remote desktop / file transfer* framing, we additionally measure user-visible, session-level metrics and quantify how one-time handshake cost amortizes over longer sessions and larger transfers.

Metrics: (i) T_{connect} (connect start $\rightarrow$ handshake established, including explicit Finished key confirmation and event emission), (ii) $T_{\text{first_frame}}$ (connect start $\rightarrow$ first decrypted remote-desktop frame processed by the rendering path), and (iii) T_{total} (connect start $\rightarrow$ completion of a bulk transfer over the negotiated session keys).

Method: We reuse the production `HandshakeDriver` and negotiated `SessionKeys`, with the following setup:

- Workload: one encrypted “first frame” plus bulk transfer workloads (1–100 MiB).




Fig. 10. System-level impact and amortization (p95). Top: T_{connect} across RTT+jitter conditions. Bottom: total session time ($T_{\text{total}} = \text{connect} + \text{transfer}$) versus file size under RTT 50±20 ms, showing that one-time handshake overhead is amortized by larger transfers. Data from `Artifacts/system_impact_2026-01-23.csv`.

- Channel model: deterministic RTT+jitter on a reliable stream.
- Throughput control: bulk payloads are rate-limited (MiB/s) to avoid per-chunk RTT artifacts.
- Output: date-stamped system-impact CSV artifact generated by `SystemImpactBenchTests`.

Results: Supplementary Table S11 reports p50/p95 T_{connect} and T_{total} results under the controlled RTT+jitter profiles.

Traffic Analysis Mitigation (SBP2 Padding). To reduce attempt identifiability on the wire, SkyBridge optionally applies a constant-size framing layer (SBP2) that buckets framed payloads into powers-of-two-sized envelopes. Here, **SBP1** denotes *handshake-frame padding* (MessageA/MessageB/Finished are padded to fixed buckets before transport), while **SBP2** denotes *traffic framing and padding* applied to all framed payloads (control and encrypted data-plane) with a per-frame header and bucketed padding up to a configured cap. We measure SBP2 behavior on a real protocol trace: an end-to-end handshake (SBP1 + SBP2 enabled) on an in-memory loopback transport, followed by bidirectional AES-GCM-encrypted data-plane frames derived from the negotiated session keys. Supplementary Table S15 reports




Fig. 11. SBP2 padding overhead across representative handshake, control, and data-plane workloads (macOS 26.x). Overhead is computed as $(\text{paddedBytes}/\text{rawBytes} - 1) \times 100\%$ aggregated per label over a deterministic trace ($N=1000$ handshake iterations; data-plane workload uses encrypted control messages and large-frame mixes). Generated from `Artifacts/traffic_padding_2026-01-23.csv`.




Fig. 12. SBP2 bucket-cap sensitivity on macOS 26.x. Lower over-cap rate means broader padding coverage across 64 KiB, 128 KiB, and 256 KiB caps for remote-desktop and file-transfer mixes.

the resulting quantization behavior and overhead; the full per-label bucket histograms are included in `Artifacts/traffic_padding_2026-01-23.csv`.

Sensitivity Study (Bucket Cap). We vary the maximum bucket size for SBP2 ($cap = 64 \text{ KiB}, 128 \text{ KiB}, 256 \text{ KiB}$) and re-run the same deterministic protocol trace. We report two overhead definitions: *all-frames overhead* (includes passthrough frames) and *padded-only overhead* (restricted to frames that entered bucket padding), alongside *over-cap rate* (fraction of frames whose framed payload exceeds cap) and *bucket entropy* (Shannon entropy of padded bucket sizes per label). When over-cap is high, all-frames overhead may approach zero while padded-only overhead remains non-zero; this is a cap tradeoff, not a measurement bug. Supplementary Table S16 summarizes this overhead–privacy tradeoff and distinguishes remote-desktop-like vs file-chunk-like large-frame distributions.

Operational decision guide (rule-of-thumb). For interactive remote-desktop-first workloads, prefer SBP2 disabled or 64 KiB cap. For mixed control + moderate transfer workloads, 128 KiB is a balanced default. For bulk-transfer-heavy, privacy-prioritized workloads, 256 KiB improves padding coverage at the cost of higher bandwidth overhead.

AEAD selection (v1). Handshake payloads and data-plane traffic both use AES-256-GCM in v1; Supplementary Table S7 summarizes the handshake-negotiated vs fixed data-plane mapping.

7.7 Network Impairments: Controlled, Kernel-level, and Real-network Evidence

To reduce external-validity ambiguity, we report a three-layer evidence stack with one shared profile source (`Scripts/netem_profiles.json`): (i) controlled in-process simulation, (ii) kernel-level emulation cross-check (`dummysnet/netem`), and (iii) real-network micro-study.

Controlled simulation (in-process). We simulate four network conditions with a retry-aware model (max 2 retries) to isolate protocol logic under reproducible loss/jitter/reorder perturbations. Loss is Bernoulli per packet; jitter is sampled uniformly from $\pm J \text{ ms}$ around base RTT; reordering is modeled as added random delay (50–150 ms) on 10% of packets.

Controlled-simulation summary: completion is ≥ 0.9989 (min), p95 latency spans 251–1033 ms, and $E_{\text{handshakeFailed}} \leq 11$ per 10k runs across the reported profiles/suites. Supplementary Table S8 reports the full controlled-simulation table.

Kernel-level cross-check.

- Tools: `dnetcl/pfctl` (macOS `dummysnet`) and `tc netem` (Linux), with aligned profile IDs.
- Orchestration: kernel emulation script; outputs are summarized in Supplementary Table S19.
- Capability note: reorder is reported only for `netem`; `dummysnet` reorder cells are *n/a*.
- Interpretation: on some macOS setups, `dummysnet` with localhost endpoints may under-shape absolute delay; therefore, delay-scale reading is anchored by `netem` and controlled simulation, while cross-tool completion robustness remains the primary check.

Real-network micro-study.

- Metrics: STUN path telemetry and TCP end-to-end timing for payload-only handshake baselines (classic 687 B; core-gate PQC 12002 B for liboqs and 12077 B for liboqs v2 FS; Apple contrast 12009 B where available), reported in Supplementary Table S17.
- Artifact source: `Artifacts/realnet_*`.
- Direct label in this snapshot: `wifi_direct` (remote EC2 over TCP 443).
- Overlay labels in this snapshot: `ec2_ssh_tunnel` and `hotspot_ssh_tunnel`.
- Reporting rule: direct and overlay paths are reported separately to distinguish external direct-path evidence from tunnel-mediated behavior under restrictive inbound NAT/firewall conditions.
- Accounting: the harness sends `4B length-prefix + payload`; payload-byte labels exclude the fixed prefix.
- External check: dedicated direct-path EC2 Linux run (no tunnel, 54.92.79.99:44444, $n=50$ per payload) in Supplementary Table S18.
- Cross-platform matrix: measured Ubuntu direct/tunnel rows plus static Android/Ubuntu contract-consistency pair rows are summarized in Supplementary Table S23.

8 LIMITATIONS AND FUTURE WORK

8.1 Current Limitations

- 1) **Platform scope:** Main performance numbers are Apple-centric; non-Apple deployments require remeasurement under the same artifact and workload controls.
- 2) **Secrecy scope:** v1 static-recipient KEM deployments do not provide the same secrecy semantics as the v2 hybrid path (Supplementary Table S22).
- 3) **External validity:** Real-network evidence remains small-scale and does not yet cover broad mobility and ISP diversity.

8.2 Future Directions

Priority extensions are broader Android/Ubuntu replication, ephemeral-ephemeral secrecy designs beyond v2, and cross-peer audit attestation.

9 CONCLUSION

SkyBridge Compass contributes an auditable migration contract for crypto-agile P2P handshakes with explicit fallback semantics, deterministic failure outcomes, and replayable evidence. In our Apple deployment instance, default policy preserves availability with local-only audited fallback, while strict policy remains fail-closed (or bootstrap-assisted with mandatory PQC rekey). These are deployment-instance measurements and should be re-established on non-Apple targets.

Acknowledgment: This work received no external funding.

DATA AND ARTIFACT AVAILABILITY

Artifacts are available at github.com/billlza/Skybridge-Compass (commit 7df9344ca0b0); frozen source archives and checksums are documented in the Supplementary Materials.

Conflict of Interest: The author declares no competing interests.

REFERENCES

- [1] Bluetooth SIG, “Bluetooth Core Specification v5.4,” 2023.
- [2] Apple Inc., “Continuity,” Apple Platform Security Guide, 2024.
- [3] T. Taubert and C. A. Wood, “SPAKE2+, an Augmented PAKE,” RFC 9383, IETF, Sept. 2023.
- [4] E. Rescorla, “The Transport Layer Security (TLS) Protocol Version 1.3,” RFC 8446, IETF, 2018.
- [5] J. Iyengar and M. Thomson, “QUIC: A UDP-Based Multiplexed and Secure Transport,” RFC 9000, IETF, 2021.
- [6] E. Barker and A. Roginsky, “Transitioning the Use of Cryptographic Algorithms and Key Lengths,” NIST SP 800-131A Rev. 2, 2019.
- [7] E. Barker et al., “Considerations for Achieving Cryptographic Agility: Strategies and Practices,” NIST Cybersecurity White Paper (CSWP) 39, Dec. 2025, doi: 10.6028/NIST.CSWP.39.
- [8] T. Perrin, “The Noise Protocol Framework,” 2018.
- [9] NIST, “Post-Quantum Cryptography Standardization,” 2024. Accessed: Dec. 2025.
- [10] Signal Foundation, “PQXDH Key Agreement Protocol,” Signal Technical Documentation, 2023.
- [11] Cloudflare, “Post-Quantum Cryptography Goes GA,” Sept. 2023. Accessed: Dec. 2025.
- [12] P. Schwabe, D. Stebila, and T. Wiggers, “Post-Quantum TLS Without Handshake Signatures,” in Proc. ACM CCS 2020, pp. 1461–1480.
- [13] E. Alkim, L. Ducas, T. Pöppelmann, and P. Schwabe, “Post-Quantum Key Exchange—A New Hope,” in Proc. USENIX Security 2016, pp. 327–343.
- [14] J. Bos et al., “CRYSTALS—Kyber: A CCA-Secure Module-Lattice-Based KEM,” in Proc. IEEE EuroS&P 2018, pp. 353–367.
- [15] L. Ducas et al., “CRYSTALS—Dilithium: A Lattice-Based Digital Signature Scheme,” IACR Trans. CHES, vol. 2018, no. 1, pp. 238–268.
- [16] Apple Inc., “Quantum-secure cryptography in Apple operating systems,” Apple Platform Security, Jan. 2026. Accessed: Feb. 2026.
- [17] D. Beyer et al., “Software Model Checking,” in Handbook of Model Checking, Springer, 2018.
- [18] B. Beurdouche et al., “A Messy State of the Union: Taming the Composite State Machines of TLS,” in Proc. IEEE S&P 2015, pp. 535–552.
- [19] Apple Inc., “Swift Concurrency,” The Swift Programming Language, 2024.
- [20] E. Rescorla and N. Modadugu, “Datagram Transport Layer Security Version 1.2,” RFC 6347, IETF, Jan. 2012.
- [21] M. Green and M. Smith, “Developers Are Not the Enemy!: The Need for Usable Security APIs,” IEEE Security & Privacy, vol. 14, no. 5, pp. 40–46, Sept./Oct. 2016.
- [22] K. Bhargavan, C. Fournet, M. Kohlweiss, A. Pironti, and P.-Y. Strub, “Implementing TLS with Verified Cryptographic Security,” in Proc. IEEE S&P 2013, pp. 445–459.
- [23] Apple Inc., “CryptoKit Framework,” Apple Developer Documentation, 2025. Accessed: Dec. 2025.
- [24] NIST, “Module-Lattice-Based Digital Signature Standard,” FIPS 204, Aug. 2024.
- [25] D. McGrew, “Achieving Crypto Agility,” in Proc. RSA Conference, 2019.
- [26] R. Barnes, K. Bhargavan, B. Lipp, and C. Wood, “Hybrid Public Key Encryption,” RFC 9180, IETF, Feb. 2022.
- [27] Apple Inc., “Get ahead with quantum-secure cryptography,” WWDC 2025 Session 314, June 2025. Accessed: Dec. 2025.
- [28] NIST, “Module-Lattice-Based Key-Encapsulation Mechanism Standard,” FIPS 203, Aug. 2024.
- [29] NIST, “Recommendations for Key-Encapsulation Mechanisms,” SP 800-227, Final, Nov. 2025.
- [30] K. Cremers, M. Horvat, J. Hoyland, and S. Scott, “A Comprehensive Symbolic Analysis of TLS 1.3,” IACR ePrint 2020/1044, 2020.
- [31] P. Schwabe, D. Stebila, and T. Wiggers, “KEMTLS: Authenticated Key Exchange without Signatures in TLS 1.3,” 2020.
- [32] S. Meier, B. Schmidt, C. Cremers, and D. Basin, “The Tamarin Prover for the Symbolic Analysis of Security Protocols,” in Proc. CAV 2013, LNCS 8044, pp. 696–701.
- [33] M. Barbosa, D. Connolly, J. Duarte, A. Kaiser, P. Schwabe, K. Varber, and B. Westerbaan, “X-Wing: The Hybrid KEM You’ve Been Looking For,” IETF Internet-Draft, draft-connolly-cfrg-xwing-kem-09, Jan. 2026.
- [34] D. J. Bernstein, “Curve25519: New Diffie-Hellman Speed Records,” in Proc. PKC 2006, LNCS 3958, Springer, 2006, pp. 207–228.
- [35] D. J. Bernstein, N. Duif, T. Lange, P. Schwabe, and B.-Y. Yang, “High-Speed High-Security Signatures,” J. Cryptographic Engineering, vol. 2, no. 2, pp. 77–89, 2012.
- [36] H. Krawczyk, “SIGMA: The ‘SIGn-and-MAC’ Approach to Authenticated Diffie-Hellman and Its Use in the IKE Protocols,” in Proc. CRYPTO 2003, LNCS 2729, pp. 400–425.
- [37] H. Krawczyk and P. Eronen, “HMAC-based Extract-and-Expand Key Derivation Function (HKDF),” RFC 5869, IETF, May 2010.
- [38] M. Dworkin, “Recommendation for Block Cipher Modes of Operation: Galois/Counter Mode (GCM) and GMAC,” NIST SP 800-38D, Nov. 2007.
- [39] N. Aviram et al., “DROWN: Breaking TLS Using SSLv2,” in Proc. USENIX Security 2016, pp. 689–706.
- [40] D. Adrian et al., “Imperfect Forward Secrecy: How Diffie-Hellman Fails in Practice,” in Proc. ACM CCS 2015, pp. 5–17.
- [41] C. Meyer and J. Schwenk, “Lessons Learned From Previous SSL/TLS Attacks—A Brief Chronology of Attacks and Weaknesses,” IACR Cryptology ePrint Archive, Report 2013/049, 2013.